

Chapter 4.1 - 4.6

Chapter 4 The tidyverse

Up to now we have been manipulating vectors by reordering and subsetting them through indexing. However, once we start more advanced analyses, the preferred unit for data storage is not the vector but the data frame. In this chapter we learn to work directly with data frames, which greatly facilitate the organization of information. We will be using data frames for the majority of this book. We will focus on a specific data format referred to as tidy and on specific collection of packages that are particularly helpful for working with tidy data referred to as the tidyverse.

We can load all the tidyverse packages at once by installing and loading the tidyverse package:

```
library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.1 --

## v ggplot2 3.3.5      v purrr  0.3.4
## v tibble  3.1.4      v dplyr  1.0.7
## v tidyr   1.1.3      v stringr 1.4.0
## v readr   2.0.1      v forcats 0.5.1

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
```

We will learn how to implement the tidyverse approach throughout the book, but before delving into the details, in this chapter we introduce some of the most widely used tidyverse functionality, starting with the dplyr package for manipulating data frames and the purrr package for working with functions. Note that the tidyverse also includes a graphing package, ggplot2, which we introduce later in Chapter 7 in the Data Visualization part of the book; the readr package discussed in Chapter 5; and many others. In this chapter, we first introduce the concept of tidy data and then demonstrate how we use the tidyverse to work with data frames in this format.

4.1 Tidy data

We say that a data table is in tidy format if each row represents one observation and columns represent the different variables available for each of these observations. The murders dataset is an example of a tidy data frame.

```
#>      state abb region population total
#> 1  Alabama  AL  South    4779736   135
#> 2  Alaska   AK   West     710231    19
#> 3  Arizona  AZ   West    6392017   232
#> 4  Arkansas AR  South    2915918    93
#> 5 California CA  West    37253956  1257
#> 6  Colorado CO  West     5029196    65
```

Each row represent a state with each of the five columns providing a different variable related to these states: name, abbreviation, region, population, and total murders.

To see how the same information can be provided in different formats, consider the following example:

```
#>      country year fertility
#> 1    Germany 1960      2.41
#> 2 South Korea 1960      6.16
#> 3    Germany 1961      2.44
#> 4 South Korea 1961      5.99
#> 5    Germany 1962      2.47
#> 6 South Korea 1962      5.79
```

This tidy dataset provides fertility rates for two countries across the years. This is a tidy dataset because each row presents one observation with the three variables being country, year, and fertility rate. However, this dataset originally came in another format and was reshaped for the dslabs package. Originally, the data was in the following format:

```
#>      country 1960 1961 1962
#> 1    Germany 2.41 2.44 2.47
#> 2 South Korea 6.16 5.99 5.79
```

The same information is provided, but there are two important differences in the format: 1) each row includes several observations and 2) one of the variables, year, is stored in the header. For the tidyverse packages to be optimally used, data need to be reshaped into tidy format, which you will learn to do in the Data Wrangling part of the book. Until then, we will use example datasets that are already in tidy format.

Although not immediately obvious, as you go through the book you will start to appreciate the advantages of working in a framework in which functions use tidy formats for both inputs and outputs. You will see how this permits the data analyst to focus on more important aspects of the analysis rather than the format of the data.

4.2 Exercises

1. Examine the built-in dataset `co2`. Which of the following is true:
 - a. `co2` is tidy data: it has one year for each row.
 - b. `co2` is not tidy: we need at least one column with a character vector.
 - c. `co2` is not tidy: it is a matrix instead of a data frame.
 - d. **`co2` is not tidy: to be tidy we would have to wrangle it to have three columns (year, month and value), then each `co2` observation would have a row.**

```
head(co2)
```

```
## [1] 315.42 316.31 316.50 317.56 318.13 318.00
```

2. Examine the built-in dataset `ChickWeight`. Which of the following is true:
 - a. `ChickWeight` is not tidy: each chick has more than one row.
 - b. `ChickWeight` is tidy: each observation (a weight) is represented by one row. The chick from which this measurement came is one of the variables.
 - c. `ChickWeight` is not tidy: we are missing the year column.
 - d. **`ChickWeight` is tidy: it is stored in a data frame.**

```
head(ChickWeight)
```

```
##   weight Time Chick Diet
## 1     42   0     1     1
## 2     51   2     1     1
## 3     59   4     1     1
## 4     64   6     1     1
## 5     76   8     1     1
## 6     93  10     1     1
```

3. Examine the built-in dataset BOD. Which of the following is true:
- BOD is not tidy: it only has six rows.
 - BOD is not tidy: the first column is just an index.
 - BOD is tidy: each row is an observation with two values (time and demand)**
 - BOD is tidy: all small datasets are tidy by definition.

```
head(BOD)
```

```
##   Time demand
## 1     1    8.3
## 2     2   10.3
## 3     3   19.0
## 4     4   16.0
## 5     5   15.6
## 6     7   19.8
```

4. Which of the following built-in datasets is tidy (you can pick more than one):
- BJsales
 - EuStockMarkets
 - DNase**
 - Formaldehyde**
 - Orange**
 - UCBAdmissions

```
head(DNase)
```

```
##   Run      conc density
## 1   1 0.04882812  0.017
## 2   1 0.04882812  0.018
## 3   1 0.19531250  0.121
## 4   1 0.19531250  0.124
## 5   1 0.39062500  0.206
## 6   1 0.39062500  0.215
```

```
head(Formaldehyde)
```

```
## carb optden
## 1 0.1 0.086
## 2 0.3 0.269
## 3 0.5 0.446
## 4 0.6 0.538
## 5 0.7 0.626
## 6 0.9 0.782
```

```
head(Orange)
```

```
## Tree age circumference
## 1 1 118 30
## 2 1 484 58
## 3 1 664 87
## 4 1 1004 115
## 5 1 1231 120
## 6 1 1372 142
```

4.3 Manipulating data frames

The dplyr package from the tidyverse introduces functions that perform some of the most common operations when working with data frames and uses names for these functions that are relatively easy to remember. For instance, to change the data table by adding a new column, we use mutate. To filter the data table to a subset of rows, we use filter. Finally, to subset the data by selecting specific columns, we use select.

4.3.1 Adding a column with mutate

We want all the necessary information for our analysis to be included in the data table. So the first task is to add the murder rates to our murders data frame. The function mutate takes the data frame as a first argument and the name and values of the variable as a second argument using the convention name = values. So, to add murder rates, we use:

```
library(dslabs)
data("murders")
murders <- mutate(murders, rate = total / population * 100000)
```

Notice that here we used total and population inside the function, which are objects that are not defined in our workspace. But why don't we get an error?

This is one of dplyr's main features. Functions in this package, such as mutate, know to look for variables in the data frame provided in the first argument. In the call to mutate above, total will have the values in murders\$total. This approach makes the code much more readable.

We can see that the new column is added:

```
head(murders)
```

```
## state abb region population total rate
## 1 Alabama AL South 4779736 135 2.824424
## 2 Alaska AK West 710231 19 2.675186
## 3 Arizona AZ West 6392017 232 3.629527
## 4 Arkansas AR South 2915918 93 3.189390
## 5 California CA West 37253956 1257 3.374138
## 6 Colorado CO West 5029196 65 1.292453
```

Although we have overwritten the original murders object, this does not change the object that loaded with `data(murders)`. If we load the murders data again, the original will overwrite our mutated version.

4.3.2 Subsetting with filter

Now suppose that we want to filter the data table to only show the entries for which the murder rate is lower than 0.71. To do this we use the filter function, which takes the data table as the first argument and then the conditional statement as the second. Like mutate, we can use the unquoted variable names from murders inside the function and it will know we mean the columns and not objects in the workspace.

```
filter(murders, rate <= 0.71)
```

##	state	abb	region	population	total	rate
## 1	Hawaii	HI	West	1360301	7	0.5145920
## 2	Iowa	IA	North Central	3046355	21	0.6893484
## 3	New Hampshire	NH	Northeast	1316470	5	0.3798036
## 4	North Dakota	ND	North Central	672591	4	0.5947151
## 5	Vermont	VT	Northeast	625741	2	0.3196211

4.3.3 Selecting columns with select

Although our data table only has six columns, some data tables include hundreds. If we want to view just a few, we can use the dplyr select function. In the code below we select three columns, assign this to a new object and then filter the new object:

```
new_table <- select(murders, state, region, rate)
filter(new_table, rate <= 0.71)
```

##	state	region	rate
## 1	Hawaii	West	0.5145920
## 2	Iowa	North Central	0.6893484
## 3	New Hampshire	Northeast	0.3798036
## 4	North Dakota	North Central	0.5947151
## 5	Vermont	Northeast	0.3196211

In the call to select, the first argument murders is an object, but state, region, and rate are variable names.

Exercises

1. Load the dplyr package and the murders dataset.

```
library(dplyr)
library(dslabs)
data(murders)
```

You can add columns using the dplyr function mutate. This function is aware of the column names and inside the function you can call them unquoted:

```
murders <- mutate(murders, population_in_millions = population / 10^6)
```

Use the function `mutate` to add a `murders` column named `rate` with the per 100,000 murder rate as in the example code above. Make sure you redefine `murders` as done in the example code above (`murders <- [your code]`) so we can keep using this variable.

```
murders <- mutate(murders, rate = total / population * 100000)
murders
```

	state	abb	region	population	total
## 1	Alabama	AL	South	4779736	135
## 2	Alaska	AK	West	710231	19
## 3	Arizona	AZ	West	6392017	232
## 4	Arkansas	AR	South	2915918	93
## 5	California	CA	West	37253956	1257
## 6	Colorado	CO	West	5029196	65
## 7	Connecticut	CT	Northeast	3574097	97
## 8	Delaware	DE	South	897934	38
## 9	District of Columbia	DC	South	601723	99
## 10	Florida	FL	South	19687653	669
## 11	Georgia	GA	South	9920000	376
## 12	Hawaii	HI	West	1360301	7
## 13	Idaho	ID	West	1567582	12
## 14	Illinois	IL	North Central	12830632	364
## 15	Indiana	IN	North Central	6483802	142
## 16	Iowa	IA	North Central	3046355	21
## 17	Kansas	KS	North Central	2853118	63
## 18	Kentucky	KY	South	4339367	116
## 19	Louisiana	LA	South	4533372	351
## 20	Maine	ME	Northeast	1328361	11
## 21	Maryland	MD	South	5773552	293
## 22	Massachusetts	MA	Northeast	6547629	118
## 23	Michigan	MI	North Central	9883640	413
## 24	Minnesota	MN	North Central	5303925	53
## 25	Mississippi	MS	South	2967297	120
## 26	Missouri	MO	North Central	5988927	321
## 27	Montana	MT	West	989415	12
## 28	Nebraska	NE	North Central	1826341	32
## 29	Nevada	NV	West	2700551	84
## 30	New Hampshire	NH	Northeast	1316470	5
## 31	New Jersey	NJ	Northeast	8791894	246
## 32	New Mexico	NM	West	2059179	67
## 33	New York	NY	Northeast	19378102	517
## 34	North Carolina	NC	South	9535483	286
## 35	North Dakota	ND	North Central	672591	4
## 36	Ohio	OH	North Central	11536504	310
## 37	Oklahoma	OK	South	3751351	111
## 38	Oregon	OR	West	3831074	36
## 39	Pennsylvania	PA	Northeast	12702379	457
## 40	Rhode Island	RI	Northeast	1052567	16
## 41	South Carolina	SC	South	4625364	207
## 42	South Dakota	SD	North Central	814180	8
## 43	Tennessee	TN	South	6346105	219

## 44	Texas	TX	South	25145561	805
## 45	Utah	UT	West	2763885	22
## 46	Vermont	VT	Northeast	625741	2
## 47	Virginia	VA	South	8001024	250
## 48	Washington	WA	West	6724540	93
## 49	West Virginia	WV	South	1852994	27
## 50	Wisconsin	WI	North Central	5686986	97
## 51	Wyoming	WY	West	563626	5
##	population_in_millions		rate		
## 1	4.779736		2.8244238		
## 2	0.710231		2.6751860		
## 3	6.392017		3.6295273		
## 4	2.915918		3.1893901		
## 5	37.253956		3.3741383		
## 6	5.029196		1.2924531		
## 7	3.574097		2.7139722		
## 8	0.897934		4.2319369		
## 9	0.601723		16.4527532		
## 10	19.687653		3.3980688		
## 11	9.920000		3.7903226		
## 12	1.360301		0.5145920		
## 13	1.567582		0.7655102		
## 14	12.830632		2.8369608		
## 15	6.483802		2.1900730		
## 16	3.046355		0.6893484		
## 17	2.853118		2.2081106		
## 18	4.339367		2.6732010		
## 19	4.533372		7.7425810		
## 20	1.328361		0.8280881		
## 21	5.773552		5.0748655		
## 22	6.547629		1.8021791		
## 23	9.883640		4.1786225		
## 24	5.303925		0.9992600		
## 25	2.967297		4.0440846		
## 26	5.988927		5.3598917		
## 27	0.989415		1.2128379		
## 28	1.826341		1.7521372		
## 29	2.700551		3.1104763		
## 30	1.316470		0.3798036		
## 31	8.791894		2.7980319		
## 32	2.059179		3.2537239		
## 33	19.378102		2.6679599		
## 34	9.535483		2.9993237		
## 35	0.672591		0.5947151		
## 36	11.536504		2.6871225		
## 37	3.751351		2.9589340		
## 38	3.831074		0.9396843		
## 39	12.702379		3.5977513		
## 40	1.052567		1.5200933		
## 41	4.625364		4.4753235		
## 42	0.814180		0.9825837		
## 43	6.346105		3.4509357		
## 44	25.145561		3.2013603		
## 45	2.763885		0.7959810		

```
## 46          0.625741  0.3196211
## 47          8.001024  3.1246001
## 48          6.724540  1.3829942
## 49          1.852994  1.4571013
## 50          5.686986  1.7056487
## 51          0.563626  0.8871131
```

2. If `rank(x)` gives you the ranks of `x` from lowest to highest, `rank(-x)` gives you the ranks from highest to lowest. Use the function `mutate` to add a column `rank` containing the rank, from highest to lowest murder rate. Make sure you redefine `murders` so we can keep using this variable.

```
murders <- mutate(murders, rank = rank(-rate))
murders$rank
```

```
## [1] 23 27 10 17 14 38 25  6  1 13  9 49 46 22 31 47 30 28  2 44  4 32  7 40  8
## [26]  3 39 33 19 50 24 15 29 20 48 26 21 42 11 35  5 41 12 16 45 51 18 37 36 34
## [51] 43
```

3. With `dplyr`, we can use `select` to show only certain columns. For example, with this code we would only show the states and population sizes:

```
select(murders, state, population) %>% head()
```

```
##      state population
## 1  Alabama    4779736
## 2   Alaska    710231
## 3   Arizona    6392017
## 4  Arkansas    2915918
## 5 California    37253956
## 6   Colorado    5029196
```

Use `select` to show the state names and abbreviations in `murders`. Do not redefine `murders`, just show the results.

```
select(murders, state, abb)
```

```
##      state abb
## 1  Alabama  AL
## 2   Alaska  AK
## 3   Arizona  AZ
## 4  Arkansas  AR
## 5 California  CA
## 6   Colorado  CO
## 7 Connecticut  CT
## 8   Delaware  DE
## 9 District of Columbia  DC
## 10  Florida  FL
## 11   Georgia  GA
## 12   Hawaii  HI
## 13   Idaho  ID
## 14  Illinois  IL
```



```
## 15      Indiana IN
## 16      Iowa IA
## 17      Kansas KS
## 18      Kentucky KY
## 19      Louisiana LA
## 20      Maine ME
## 21      Maryland MD
## 22      Massachusetts MA
## 23      Michigan MI
## 24      Minnesota MN
## 25      Mississippi MS
## 26      Missouri MO
## 27      Montana MT
## 28      Nebraska NE
## 29      Nevada NV
## 30      New Hampshire NH
## 31      New Jersey NJ
## 32      New Mexico NM
## 33      New York NY
## 34      North Carolina NC
## 35      North Dakota ND
## 36      Ohio OH
## 37      Oklahoma OK
## 38      Oregon OR
## 39      Pennsylvania PA
## 40      Rhode Island RI
## 41      South Carolina SC
## 42      South Dakota SD
## 43      Tennessee TN
## 44      Texas TX
## 45      Utah UT
## 46      Vermont VT
## 47      Virginia VA
## 48      Washington WA
## 49      West Virginia WV
## 50      Wisconsin WI
## 51      Wyoming WY
```

4. The `dplyr` function `filter` is used to choose specific rows of the data frame to keep. Unlike `select` which is for columns, `filter` is for rows. For example, you can show just the New York row like this:

```
filter(murders, state == "New York")
```

```
##      state abb    region population total population_in_millions    rate rank
## 1 New York  NY Northeast   19378102    517                19.3781 2.66796    29
```

You can use other logical vectors to filter rows.

Use `filter` to show the top 5 states with the highest murder rates. After we add murder rate and rank, do not change the `murders` dataset, just show the result. Remember that you can filter based on the rank column.

```
filter(murders, rank <= 5)
```

```
##           state abb      region population total
## 1 District of Columbia DC      South      601723    99
## 2      Louisiana LA      South    4533372    351
## 3      Maryland MD      South    5773552    293
## 4      Missouri MO North Central    5988927    321
## 5    South Carolina SC      South    4625364    207
## population_in_millions      rate rank
## 1           0.601723 16.452753    1
## 2           4.533372  7.742581    2
## 3           5.773552  5.074866    4
## 4           5.988927  5.359892    3
## 5           4.625364  4.475323    5
```

5. We can remove rows using the `!=` operator. For example, to remove Florida, we would do this:

```
no_florida <- filter(murders, state != "Florida")
```

Create a new data frame called `no_south` that removes states from the South region. How many states are in this category? You can use the function `nrow` for this.

```
no_south <- filter(murders, region != "South")
nrow(no_south)
```

```
## [1] 34
```

6. We can also use `%in%` to filter with `dplyr`. You can therefore see the data from New York and Texas like this:

```
filter(murders, state %in% c("New York", "Texas"))
```

```
##      state abb      region population total population_in_millions      rate rank
## 1 New York NY Northeast   19378102    517           19.37810 2.66796    29
## 2   Texas TX      South   25145561    805           25.14556 3.20136    16
```

Create a new data frame called `murders_nw` with only the states from the Northeast and the West. How many states are in this category?

```
murders_nw <- filter(murders, region %in% c("Northeast", "West"))
nrow(murders_nw)
```

```
## [1] 22
```

7. Suppose you want to live in the Northeast or West and want the murder rate to be less than 1. We want to see the data for the states satisfying these options. Note that you can use logical operators with `filter`. Here is an example in which we filter to keep only small states in the Northeast region.

```
filter(murders, population < 5000000 & region == "Northeast")
```

```
##           state abb      region population total population_in_millions      rate
## 1    Connecticut  CT Northeast   3574097     97             3.574097 2.7139722
## 2         Maine   ME Northeast   1328361    11             1.328361 0.8280881
## 3 New Hampshire  NH Northeast   1316470     5             1.316470 0.3798036
## 4  Rhode Island  RI Northeast   1052567    16             1.052567 1.5200933
## 5      Vermont   VT Northeast    625741     2             0.625741 0.3196211
##   rank
## 1    25
## 2    44
## 3    50
## 4    35
## 5    51
```

Make sure `murders` has been defined with `rate` and `rank` and still has all states. Create a table called `my_states` that contains rows for states satisfying both the conditions: it is in the Northeast or West and the murder rate is less than 1. Use `select` to show only the state name, the rate, and the rank.

```
my_states <- filter(murders, rate < 1 & region %in% c("Northeast", "West"))
select(my_states, state, rate, rank)
```

```
##           state      rate rank
## 1      Hawaii 0.5145920   49
## 2      Idaho 0.7655102   46
## 3      Maine 0.8280881   44
## 4 New Hampshire 0.3798036   50
## 5      Oregon 0.9396843   42
## 6      Utah 0.7959810   45
## 7      Vermont 0.3196211   51
## 8      Wyoming 0.8871131   43
```

4.5 The pipe: %>%

With `dplyr` we can perform a series of operations, for example `select` and then `filter`, by sending the results of one function to another using what is called the pipe operator: `%>%`. Some details are included below.

We wrote code above to show three variables (`state`, `region`, `rate`) for states that have murder rates below 0.71. To do this, we defined the intermediate object `new_table`. In `dplyr` we can write code that looks more like a description of what we want to do without intermediate objects:

original data → `select` → `filter`

For such an operation, we can use the pipe `%>%`. The code looks like this:

```
murders %>% select(state, region, rate) %>% filter(rate <= 0.71)
```

```
##           state      region      rate
## 1      Hawaii          West 0.5145920
## 2      Iowa North Central 0.6893484
## 3 New Hampshire Northeast 0.3798036
## 4 North Dakota North Central 0.5947151
## 5      Vermont Northeast 0.3196211
```

This line of code is equivalent to the two lines of code above. What is going on here?

In general, the pipe sends the result of the left side of the pipe to be the first argument of the function on the right side of the pipe. Here is a very simple example:

```
16 %>% sqrt()
```

```
## [1] 4
```

We can continue to pipe values along:

```
16 %>% sqrt() %>% log2()
```

```
## [1] 2
```

The above statement is equivalent to `log2(sqrt(16))`.

Remember that the pipe sends values to the first argument, so we can define other arguments as if the first argument is already defined:

```
16 %>% sqrt() %>% log(base = 2)
```

```
## [1] 2
```

Therefore, when using the pipe with data frames and dplyr, we no longer need to specify the required first argument since the dplyr functions we have described all take the data as the first argument. In the code we wrote:

```
murders %>% select(state, region, rate) %>% filter(rate <= 0.71)
```

```
##           state           region      rate
## 1      Hawaii           West 0.5145920
## 2      Iowa North Central 0.6893484
## 3 New Hampshire      Northeast 0.3798036
## 4 North Dakota North Central 0.5947151
## 5      Vermont      Northeast 0.3196211
```

`murders` is the first argument of the `select` function, and the new data frame (formerly `new_table`) is the first argument of the `filter` function.

Note that the pipe works well with functions where the first argument is the input data. Functions in tidyverse packages like dplyr have this format and can be used easily with the pipe.

4.6 Exercises

1. The pipe `%>%` can be used to perform operations sequentially without having to define intermediate objects. Start by redefining `murder` to include `rate` and `rank`.

```
murders <- mutate(murders, rate = total / population * 100000,
                  rank = rank(-rate))
```

In the solution to the previous exercise, we did the following:

```
my_states <- filter(murders, region %in% c("Northeast", "West") &
  rate < 1)

select(my_states, state, rate, rank)
```

```
##           state      rate rank
## 1      Hawaii 0.5145920   49
## 2       Idaho 0.7655102   46
## 3       Maine 0.8280881   44
## 4 New Hampshire 0.3798036   50
## 5       Oregon 0.9396843   42
## 6        Utah 0.7959810   45
## 7    Vermont 0.3196211   51
## 8     Wyoming 0.8871131   43
```

The pipe `%>%` permits us to perform both operations sequentially without having to define an intermediate variable `my_states`. We therefore could have mutated and selected in the same line like this:

```
mutate(murders, rate = total / population * 100000,
  rank = rank(-rate)) %>%
  select(state, rate, rank)
```

```
##           state      rate rank
## 1      Alabama 2.8244238   23
## 2      Alaska 2.6751860   27
## 3      Arizona 3.6295273   10
## 4      Arkansas 3.1893901   17
## 5      California 3.3741383   14
## 6      Colorado 1.2924531   38
## 7    Connecticut 2.7139722   25
## 8      Delaware 4.2319369    6
## 9 District of Columbia 16.4527532    1
## 10     Florida 3.3980688   13
## 11     Georgia 3.7903226    9
## 12     Hawaii 0.5145920   49
## 13     Idaho 0.7655102   46
## 14     Illinois 2.8369608   22
## 15     Indiana 2.1900730   31
## 16      Iowa 0.6893484   47
## 17      Kansas 2.2081106   30
## 18     Kentucky 2.6732010   28
## 19     Louisiana 7.7425810    2
## 20      Maine 0.8280881   44
## 21     Maryland 5.0748655    4
## 22 Massachusetts 1.8021791   32
## 23      Michigan 4.1786225    7
## 24      Minnesota 0.9992600   40
## 25     Mississippi 4.0440846    8
## 26      Missouri 5.3598917    3
## 27      Montana 1.2128379   39
## 28      Nebraska 1.7521372   33
## 29      Nevada 3.1104763   19
```

## 30	New Hampshire	0.3798036	50
## 31	New Jersey	2.7980319	24
## 32	New Mexico	3.2537239	15
## 33	New York	2.6679599	29
## 34	North Carolina	2.9993237	20
## 35	North Dakota	0.5947151	48
## 36	Ohio	2.6871225	26
## 37	Oklahoma	2.9589340	21
## 38	Oregon	0.9396843	42
## 39	Pennsylvania	3.5977513	11
## 40	Rhode Island	1.5200933	35
## 41	South Carolina	4.4753235	5
## 42	South Dakota	0.9825837	41
## 43	Tennessee	3.4509357	12
## 44	Texas	3.2013603	16
## 45	Utah	0.7959810	45
## 46	Vermont	0.3196211	51
## 47	Virginia	3.1246001	18
## 48	Washington	1.3829942	37
## 49	West Virginia	1.4571013	36
## 50	Wisconsin	1.7056487	34
## 51	Wyoming	0.8871131	43

Notice that `select` no longer has a data frame as the first argument. The first argument is assumed to be the result of the operation conducted right before the `%>%`.

Repeat the previous exercise, but now instead of creating a new object, show the result and only include the state, rate, and rank columns. Use a pipe `%>%` to do this in just one line.

```
filter(murders, region %in% c("Northeast", "West") & rate < 1) %>% select(state, rate, rank)
```

##	state	rate	rank
## 1	Hawaii	0.5145920	49
## 2	Idaho	0.7655102	46
## 3	Maine	0.8280881	44
## 4	New Hampshire	0.3798036	50
## 5	Oregon	0.9396843	42
## 6	Utah	0.7959810	45
## 7	Vermont	0.3196211	51
## 8	Wyoming	0.8871131	43

2. Reset `murders` to the original table by using `data(murders)`. Use a pipe to create a new data frame called `my_states` that considers only states in the Northeast or West which have a murder rate lower than 1, and contains only the state, rate and rank columns. The pipe should also have four components separated by three `%>%`. The code should look something like this:

```
# my_states <- murders %>%
#   mutate SOMETHING %>%
#   filter SOMETHING %>%
#   select SOMETHING
```

```
my_states <- murders %>% mutate(rate = total / population * 100000,
                                rank = rank(-rate)) %>%
  filter(region %in% c("Northeast", "West") & rate < 1) %>%
  select(state, rate, rank)
my_states
```

##	state	rate	rank
## 1	Hawaii	0.5145920	49
## 2	Idaho	0.7655102	46
## 3	Maine	0.8280881	44
## 4	New Hampshire	0.3798036	50
## 5	Oregon	0.9396843	42
## 6	Utah	0.7959810	45
## 7	Vermont	0.3196211	51
## 8	Wyoming	0.8871131	43